



Problem Statement:

PredCatch Analytics' Australian banking client's profitability and reputation are being hit by fraudulent ATM transactions. They want PredCatch to help them in reducing and if possible completely eliminating such fraudulent transactions. PredCatch believes it can do the same by building a predictive model to catch such fraudulent transactions in real time and decline them. Your job as PredCatch's Data Scientist is to build this fraud detection & prevention predictive model in the first step. If successful, in the 2nd step you will have to present your solutions and explain how it works to the client. The data has been made available to you. The challenging part of the problem is that the data contains very few fraud instances in comparison to the overall population. To give more edge to the solution they have also collected data regarding.

- **[geo_scores]** - location of the transactions
- **[Lambda_wts]** - their own proprietary index
- **[Qset_tats]** - on network turn around times (Transaction time)
- **[instance_scores]** - vulnerability qualification score (cibil scores)

Training data contains masked variables pertaining to each transaction id .

prediction target here is 'Target' .

1: Fraudulent transactions 0: Clean transactions

Importing Libraries

```
In [5]: 1 import os, sys
        2 import numpy as np
        3 import pandas as pd
        4 import matplotlib.pyplot as plt
        5 %matplotlib inline
        6 import seaborn as sns
        7 sns.set()
        8
        9 import warnings
       10 warnings.filterwarnings('ignore')
       11
       12 # Display all the columns of the dataframe
       13 pd.pandas.set_option('display.max_columns',None)
```

Importing Datasets

```
In [6]: 1 geo = pd.read_csv('Geo_scores.csv')
        2 instance = pd.read_csv('instance_scores.csv')
        3 lambdawts = pd.read_csv('Lambda_wts.csv')
        4 qset = pd.read_csv('Qset_tats.csv')
        5 test = pd.read_csv('test_share.csv')
        6 train = pd.read_csv('train.csv')
```

Dataset Information

In [7]:

```
1 datasets = [geo, instance, lambdawts, qset, test, train]
2 names = ["geo_scores", "instance_scores", "lambdawts", "qset", "test", "tra
3
4 # ANSI escape code for bold text
5 bold_text = "\033[1m"
6 # ANSI escape code to reset text formatting
7 reset_text = "\033[0m"
8
9 for name, dataset in zip(names, datasets):
10     # Print shape
11     print(f"{bold_text} {name} shape:{reset_text} {dataset.shape}")
12     print("-----" * 10)
13
14     # Print head
15     print(f"{bold_text}{name} head:{reset_text}")
16     print(dataset.head())
17     print("-----" * 10)
```

\	id	Group	Per1	Per2	Per3	Per4	Per5	Per6
0	146574	Grp229	-0.300000	1.540000	0.220000	-0.280000	0.570000	0.260000
1	268759	Grp141	0.633333	0.953333	0.810000	0.466667	0.910000	0.253333
2	59727	Grp188	1.043333	0.740000	0.860000	1.006667	0.583333	0.616667
3	151544	Grp426	1.283333	0.300000	0.576667	0.636667	0.256667	0.543333
4	155008	Grp443	1.186667	0.326667	0.476667	0.866667	0.436667	0.680000

	Per7	Per8	Per9	Dem1	Dem2	Dem3	Dem4	\
0	0.700000	1.076667	0.930000	0.156667	0.546667	0.530000	0.876667	
1	1.040000	0.550000	0.543333	0.433333	0.966667	0.760000	0.576667	
2	0.630000	0.686667	0.593333	1.250000	0.826667	0.826667	0.653333	
3	0.356667	0.663333	1.156667	1.186667	0.900000	0.433333	0.230000	

4	0.476667	0.686667	1.476667	1.213333	0.853333	0.583333	0.850000
	Dem5	Dem6	Dem7	Dem8	Dem9	Cred1	Cred2 \
0	0.450000	0.370000	0.786667	0.546667	0.313333	0.703333	0.813333
1	0.653333	0.553333	0.636667	0.770000	0.993333	0.536667	0.703333
2	0.663333	0.453333	0.626667	0.756667	0.953333	0.623333	0.753333
3	1.323333	0.403333	0.480000	0.460000	0.260000	0.800000	0.606667
4	1.090000	0.550000	0.706667	0.740000	0.823333	0.670000	0.896667

	Cred3	Cred4	Cred5	Cred6	Normalised_FNT
0	0.776667	0.796667	0.823333	0.783333	-249.7500
1	0.806667	0.630000	0.673333	0.673333	-249.8125
2	0.870000	0.596667	0.680000	0.670000	-248.1200
3	0.456667	0.320000	0.676667	0.660000	-222.9875
4	0.566667	0.546667	0.650000	0.663333	-196.2200

train shape: (227845, 28)

train head:

	id	Group	Per1	Per2	Per3	Per4	Per5	Per6 \
0	112751	Grp169	1.070000	0.580000	0.480000	0.766667	1.233333	1.993333
1	18495	Grp161	0.473333	1.206667	0.883333	1.430000	0.726667	0.626667
2	23915	Grp261	1.130000	0.143333	0.946667	0.123333	0.080000	0.836667
3	50806	Grp198	0.636667	1.090000	0.750000	0.940000	0.743333	0.346667
4	184244	Grp228	0.560000	1.013333	0.593333	0.416667	0.773333	0.460000

	Per7	Per8	Per9	Dem1	Dem2	Dem3	Dem4 \
0	0.340000	1.010000	0.863333	0.460000	0.643333	0.736667	0.756667
1	0.810000	0.783333	0.190000	0.470000	0.613333	0.883333	0.653333
2	0.056667	0.756667	0.226667	0.660000	0.730000	0.873333	0.923333
3	0.956667	0.633333	0.486667	1.096667	0.466667	0.670000	0.526667
4	0.853333	0.796667	0.516667	0.756667	0.683333	0.296667	0.780000

	Dem5	Dem6	Dem7	Dem8	Dem9	Cred1	Cred2 \
0	0.813333	0.693333	0.666667	0.680000	0.726667	0.606667	1.010000
1	0.463333	0.483333	0.583333	0.716667	0.743333	0.680000	0.690000
2	1.223333	0.686667	0.606667	0.690000	0.820000	0.600000	0.383333
3	0.783333	0.856667	0.716667	0.720000	0.900000	0.680000	0.846667
4	0.636667	0.783333	0.630000	0.603333	0.486667	0.693333	0.526667

	Cred3	Cred4	Cred5	Cred6	Normalised_FNT	Target
0	0.933333	0.603333	0.686667	0.673333	-245.7500	0
1	0.560000	0.670000	0.553333	0.653333	-248.0000	0
2	0.763333	0.670000	0.686667	0.673333	-233.1250	0
3	0.423333	0.520000	0.846667	0.760000	-249.7775	0
4	0.520000	0.716667	0.706667	0.673333	-247.5775	0

In [8]:

```
1 for name, dataset in zip(names, datasets):
2     # Print dataset info
3     print(f"* Information of {bold_text}{name}:{reset_text}")
4     print("-----" * 10)
5     dataset.info()
6     print("-----" * 10)
7
```

* Information of **geo_scores**:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1424035 entries, 0 to 1424034
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  ---
0    id          1424035 non-null  int64
1   geo_score   1352492 non-null  float64
dtypes: float64(1), int64(1)
memory usage: 21.7 MB
```

* Information of **instance_scores**:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1424035 entries, 0 to 1424034
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  ---
0    id          1424035 non-null  int64
1  instance_scores  1424035 non-null  float64
dtypes: float64(1), int64(1)
memory usage: 21.7 MB
```

* Information of **lambdawts**:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1400 entries, 0 to 1399
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   Group      1400 non-null   object
1  lambda_wt   1400 non-null   float64
dtypes: float64(1), object(1)
memory usage: 22.0+ KB
```

* Information of **qset**:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1424035 entries, 0 to 1424034
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  ---
0    id          1424035 non-null  int64
1  qsets_normalized_tat  1320834 non-null  float64
dtypes: float64(1), int64(1)
memory usage: 21.7 MB
```

* Information of **test**:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 56962 entries, 0 to 56961
Data columns (total 27 columns):
#   Column      Non-Null Count  Dtype
---  ---
0    id          56962 non-null  int64
1   Group      56962 non-null  object
2   Per1        56962 non-null  float64
3   Per2        56962 non-null  float64
4   Per3        56962 non-null  float64
5   Per4        56962 non-null  float64
6   Per5        56962 non-null  float64
7   Per6        56962 non-null  float64
8   Per7        56962 non-null  float64
9   Per8        56962 non-null  float64
```

```

10  Per9          56962 non-null float64
11  Dem1          56962 non-null float64
12  Dem2          56962 non-null float64
13  Dem3          56962 non-null float64
14  Dem4          56962 non-null float64
15  Dem5          56962 non-null float64
16  Dem6          56962 non-null float64
17  Dem7          56962 non-null float64
18  Dem8          56962 non-null float64
19  Dem9          56962 non-null float64
20  Cred1         56962 non-null float64
21  Cred2         56962 non-null float64
22  Cred3         56962 non-null float64
23  Cred4         56962 non-null float64
24  Cred5         56962 non-null float64
25  Cred6         56962 non-null float64
26  Normalised_FNT 56962 non-null float64
dtypes: float64(25), int64(1), object(1)
memory usage: 11.7+ MB

```

* Information of **train**:

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 227845 entries, 0 to 227844
Data columns (total 28 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   id                    227845 non-null int64
1   Group                227845 non-null object
2   Per1                 227845 non-null float64
3   Per2                 227845 non-null float64
4   Per3                 227845 non-null float64
5   Per4                 227845 non-null float64
6   Per5                 227845 non-null float64
7   Per6                 227845 non-null float64
8   Per7                 227845 non-null float64
9   Per8                 227845 non-null float64
10  Per9                 227845 non-null float64
11  Dem1                 227845 non-null float64
12  Dem2                 227845 non-null float64
13  Dem3                 227845 non-null float64
14  Dem4                 227845 non-null float64
15  Dem5                 227845 non-null float64
16  Dem6                 227845 non-null float64
17  Dem7                 227845 non-null float64
18  Dem8                 227845 non-null float64
19  Dem9                 227845 non-null float64
20  Cred1                227845 non-null float64
21  Cred2                227845 non-null float64
22  Cred3                227845 non-null float64
23  Cred4                227845 non-null float64
24  Cred5                227845 non-null float64
25  Cred6                227845 non-null float64
26  Normalised_FNT       227845 non-null float64
27  Target               227845 non-null int64
dtypes: float64(25), int64(2), object(1)
memory usage: 48.7+ MB

```

Data Preprocessing

- Checking for Duplicates

```
In [9]: 1 for name, dataset in zip(names, datasets):  
2         print(f"* Duplicated values in {bold_text}{name}:{reset_text}", dataset)  
3         print("-----" * 10)
```

* Duplicated values in **geo_scores**: 55349

* Duplicated values in **instance_scores**: 33600

* Duplicated values in **lambdawts**: 0

* Duplicated values in **qset**: 59311

* Duplicated values in **test**: 0

* Duplicated values in **train**: 0

- Checking for Missing Values

```
In [10]: 1 for name, dataset in zip(names, datasets):
          2     missing_values = dataset.isnull().sum()
          3
          4     print(f"* Details of Missing values in {bold_text}{name}:{reset_text}")
          5     print("-----" * 10)
          6     print(missing_values if not missing_values.empty else "No missing values")
          7     print("-----" * 10)
          8
```

* Details of Missing values in **geo_scores**:

id 0
geo_score 71543
dtype: int64

* Details of Missing values in **instance_scores**:

id 0
instance_scores 0
dtype: int64

* Details of Missing values in **lambdawts**:

Group 0
lambda_wt 0
dtype: int64

* Details of Missing values in **qset**:

id 0
qsets_normalized_tat 103201
dtype: int64

* Details of Missing values in **test**:

id 0
Group 0
Per1 0
Per2 0
Per3 0
Per4 0
Per5 0
Per6 0
Per7 0
Per8 0
Per9 0
Dem1 0
Dem2 0
Dem3 0
Dem4 0
Dem5 0
Dem6 0
Dem7 0
Dem8 0
Dem9 0
Cred1 0
Cred2 0
Cred3 0
Cred4 0
Cred5 0
Cred6 0
Normalised_FNT 0
dtype: int64

* Details of Missing values in **train**:

```
-----  
id                0  
Group             0  
Per1              0  
Per2              0  
Per3              0  
Per4              0  
Per5              0  
Per6              0  
Per7              0  
Per8              0  
Per9              0  
Dem1              0  
Dem2              0  
Dem3              0  
Dem4              0  
Dem5              0  
Dem6              0  
Dem7              0  
Dem8              0  
Dem9              0  
Cred1             0  
Cred2             0  
Cred3             0  
Cred4             0  
Cred5             0  
Cred6             0  
Normalised_FNT    0  
Target            0  
dtype: int64  
-----
```

```
In [11]: 1 # checking the missing values percentages  
         2  
         3 geo.isnull().sum()/len(geo)*100
```


```
Out[11]: id                0.000000  
         geo_score        5.023964  
         dtype: float64
```

```
In [12]: 1 geo.describe()
```

```
Out[12]:
```

	id	geo_score
count	1.424035e+06	1.352492e+06
mean	1.424030e+05	-9.279168e-06
std	8.221673e+04	7.827199e+00
min	0.000000e+00	-1.093900e+02
25%	7.120100e+04	-5.860000e+00
50%	1.424030e+05	1.800000e-01
75%	2.136050e+05	5.860000e+00
max	2.848060e+05	4.581000e+01

```
In [13]: 1 sns.boxplot(y='geo_score', data=geo)
          2 plt.grid();
```



```
In [14]: 1 geo['geo_score'] = geo['geo_score'].fillna(geo['geo_score'].median())
```

```
In [15]: 1 geo['geo_score'].isnull().sum()
```

```
Out[15]: 0
```

```
In [16]: 1 qset.isnull().sum()/len(qset)*100
```


```
Out[16]: id          0.000000
qsets_normalized_tat  7.247083
dtype: float64
```

```
In [17]: 1 qset.describe()
```

```
Out[17]:
```

	id	qsets_normalized_tat
count	1.424035e+06	1.320834e+06
mean	1.424030e+05	1.094006e-05
std	8.221673e+04	7.731794e+00
min	0.000000e+00	-1.404400e+02
25%	7.120100e+04	-5.860000e+00
50%	1.424030e+05	2.000000e-02
75%	2.136050e+05	5.860000e+00
max	2.848060e+05	6.110000e+01

```
In [18]: 1 sns.boxplot(y='qsets_normalized_tat', data=qset)
2 plt.grid();
```



```
In [19]: 1 qset['qsets_normalized_tat'] = qset['qsets_normalized_tat'].fillna(qset['qsets_normalized_tat'].median())
```

```
In [20]: 1 qset['qsets_normalized_tat'].isnull().sum()
```

```
Out[20]: 0
```

- Checking for duplicated values

```
In [21]: 1 for name, dataset in zip(names, datasets):
2     duplicate_values = dataset.duplicated().sum()
3
4     print(f"* Details of Duplicated values in {bold_text}{name}:{reset_text}{duplicate_values}")
5     print("-----" * 10)
```

```
* Details of Duplicated values in geo_scores:55349
```

```
* Details of Duplicated values in instance_scores:33600
```

```
* Details of Duplicated values in lambdawts:0
```

```
* Details of Duplicated values in qset:59314
```

```
* Details of Duplicated values in test:0
```

```
* Details of Duplicated values in train:0
```

In [22]: 1 train.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 227845 entries, 0 to 227844
Data columns (total 28 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    227845 non-null  int64
1   Group                 227845 non-null  object
2   Per1                  227845 non-null  float64
3   Per2                  227845 non-null  float64
4   Per3                  227845 non-null  float64
5   Per4                  227845 non-null  float64
6   Per5                  227845 non-null  float64
7   Per6                  227845 non-null  float64
8   Per7                  227845 non-null  float64
9   Per8                  227845 non-null  float64
10  Per9                  227845 non-null  float64
11  Dem1                  227845 non-null  float64
12  Dem2                  227845 non-null  float64
13  Dem3                  227845 non-null  float64
14  Dem4                  227845 non-null  float64
15  Dem5                  227845 non-null  float64
16  Dem6                  227845 non-null  float64
17  Dem7                  227845 non-null  float64
18  Dem8                  227845 non-null  float64
19  Dem9                  227845 non-null  float64
20  Cred1                 227845 non-null  float64
21  Cred2                 227845 non-null  float64
22  Cred3                 227845 non-null  float64
23  Cred4                 227845 non-null  float64
24  Cred5                 227845 non-null  float64
25  Cred6                 227845 non-null  float64
26  Normalised_FNT        227845 non-null  float64
27  Target                227845 non-null  int64
dtypes: float64(25), int64(2), object(1)
memory usage: 48.7+ MB
```

In [23]: 1 test.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 56962 entries, 0 to 56961
Data columns (total 27 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    56962 non-null  int64
1   Group                 56962 non-null  object
2   Per1                  56962 non-null  float64
3   Per2                  56962 non-null  float64
4   Per3                  56962 non-null  float64
5   Per4                  56962 non-null  float64
6   Per5                  56962 non-null  float64
7   Per6                  56962 non-null  float64
8   Per7                  56962 non-null  float64
9   Per8                  56962 non-null  float64
10  Per9                  56962 non-null  float64
11  Dem1                  56962 non-null  float64
12  Dem2                  56962 non-null  float64
13  Dem3                  56962 non-null  float64
14  Dem4                  56962 non-null  float64
15  Dem5                  56962 non-null  float64
16  Dem6                  56962 non-null  float64
17  Dem7                  56962 non-null  float64
18  Dem8                  56962 non-null  float64
19  Dem9                  56962 non-null  float64
20  Cred1                 56962 non-null  float64
21  Cred2                 56962 non-null  float64
22  Cred3                 56962 non-null  float64
23  Cred4                 56962 non-null  float64
24  Cred5                 56962 non-null  float64
25  Cred6                 56962 non-null  float64
26  Normalised_FNT       56962 non-null  float64
dtypes: float64(25), int64(1), object(1)
memory usage: 11.7+ MB
```



```
In [24]: 1 print("geo id :", geo['id'].nunique())
2 print("-----"*10)
3 print("instance id:", instance['id'].nunique())
4 print("-----"*10)
5 print("lambdawts Group :", lambdawts['Group'].nunique())
6 print("-----"*10)
7 print("qset id :", qset['id'].nunique())
8 print("-----"*10)
9 print("test - id :", test['id'].nunique())
10 print("-----"*10)
11 print("test - Group :", test['Group'].nunique())
12 print("-----"*10)
13 print("train - id :", train['id'].nunique())
14 print("-----"*10)
15 print("train - Group :", train['Group'].nunique())
```

geo id : 284807

instance id: 284807

lambdawts Group : 1400

qset id : 284807

test - id : 56962

test - Group : 915

train - id : 227845

train - Group : 1301

- geo_scores,instance_scores and Qset_stats have similar number of unique ids = **284807**
- Train has **227845** number of unique ids.
- Test has **56962** number of unique ids.
- combining id column of Train and Test : **227845 + 56962 = 284807** unique transactions ids.
- combining Group column of Train and Test : **915 + 1301 = 1406**

```
In [25]: 1 print(train.shape)
2 print(test.shape)
```

(227845, 28)

(56962, 27)

```
In [26]: 1 total_rows = train.shape[0] + test.shape[0]
2
3 print(f"Combining the Train and Test : {bold_text}{total_rows}{reset_text}")
```

Combining the Train and Test : **284807**

```
In [27]: 1 train['data'] = 'train'
2 test['data'] = 'test'      #to recognize the train and test data
```

```
In [28]: 1 all_data = pd.concat([train, test], axis=0) #combine the train and test data
```

```
In [29]: 1 all_data.head()
```

```
Out[29]:
```

	id	Group	Per1	Per2	Per3	Per4	Per5	Per6	Per7	Per8	Per9
0	112751	Grp169	1.070000	0.580000	0.480000	0.766667	1.233333	1.993333	0.340000	1.010000	0.863333
1	18495	Grp161	0.473333	1.206667	0.883333	1.430000	0.726667	0.626667	0.810000	0.783333	0.190000
2	23915	Grp261	1.130000	0.143333	0.946667	0.123333	0.080000	0.836667	0.056667	0.756667	0.226667
3	50806	Grp198	0.636667	1.090000	0.750000	0.940000	0.743333	0.346667	0.956667	0.633333	0.486667
4	184244	Grp228	0.560000	1.013333	0.593333	0.416667	0.773333	0.460000	0.853333	0.796667	0.516667

```
In [30]: 1 all_data.tail()
```

```
Out[30]:
```

	id	Group	Per1	Per2	Per3	Per4	Per5	Per6	Per7	Per8	Per9
56957	18333	Grp102	0.553333	1.043333	1.096667	0.686667	0.673333	0.340000	0.900000	0.643333	0.643333
56958	244207	Grp504	1.353333	0.616667	0.276667	0.783333	0.690000	0.650000	0.473333	0.670000	0.670000
56959	103277	Grp78	1.083333	0.433333	0.806667	0.490000	0.243333	0.316667	0.533333	0.606667	0.606667
56960	273294	Grp134	0.566667	1.153333	0.370000	0.616667	0.793333	0.226667	0.910000	0.696667	0.696667
56961	223337	Grp18	1.426667	0.110000	-0.006667	-0.200000	0.983333	1.870000	0.033333	0.963333	0.963333

```
In [31]: 1 all_data.isnull().sum()
```

```
Out[31]: id          0
Group          0
Per1           0
Per2           0
Per3           0
Per4           0
Per5           0
Per6           0
Per7           0
Per8           0
Per9           0
Dem1           0
Dem2           0
Dem3           0
Dem4           0
Dem5           0
Dem6           0
Dem7           0
Dem8           0
Dem9           0
Cred1          0
Cred2          0
Cred3          0
Cred4          0
Cred5          0
Cred6          0
Normalised_FNT 0
Target        56962
data           0
dtype: int64
```

```
In [32]: 1 all_data.shape
```

```
Out[32]: (284807, 29)
```

- Joining all the Datasets one by one who is having **same id** and **same Group**

```
In [33]: 1 geo.shape
```

```
Out[33]: (1424035, 2)
```

```
In [34]: 1 geo.head()
```

```
Out[34]:
```

	id	geo_score
0	26674	4.48
1	204314	4.48
2	176521	5.17
3	48812	-2.41
4	126870	6.55

```
In [35]: 1 geo = geo.groupby('id').mean()
```

```
In [36]: 1 all_data = pd.merge(all_data, geo, on='id', how='left')
```

```
In [37]: 1 all_data.head()
```

```
Out[37]:
```

	id	Group	Per1	Per2	Per3	Per4	Per5	Per6	Per7	Per8	Per9
0	112751	Grp169	1.070000	0.580000	0.480000	0.766667	1.233333	1.993333	0.340000	1.010000	0.863333
1	18495	Grp161	0.473333	1.206667	0.883333	1.430000	0.726667	0.626667	0.810000	0.783333	0.190000
2	23915	Grp261	1.130000	0.143333	0.946667	0.123333	0.080000	0.836667	0.056667	0.756667	0.226667
3	50806	Grp198	0.636667	1.090000	0.750000	0.940000	0.743333	0.346667	0.956667	0.633333	0.486667
4	184244	Grp228	0.560000	1.013333	0.593333	0.416667	0.773333	0.460000	0.853333	0.796667	0.516667

```
In [38]: 1 all_data.shape
```

```
Out[38]: (284807, 30)
```

```
In [39]: 1 instance['id'].nunique()
```

```
Out[39]: 284807
```

```
In [40]: 1 instance.shape
```

```
Out[40]: (1424035, 2)
```

```
In [41]: 1 instance = instance.groupby('id').mean()
```

```
In [42]: 1 instance.shape
```

```
Out[42]: (284807, 1)
```

```
In [43]: 1 all_data = pd.merge(all_data, instance, on='id', how='left')
```

```
In [44]: 1 all_data.shape
```

```
Out[44]: (284807, 31)
```

```
In [45]: 1 all_data.head()
```

```
Out[45]:
```

	id	Group	Per1	Per2	Per3	Per4	Per5	Per6	Per7	Per8	Per9
0	112751	Grp169	1.070000	0.580000	0.480000	0.766667	1.233333	1.993333	0.340000	1.010000	0.863333
1	18495	Grp161	0.473333	1.206667	0.883333	1.430000	0.726667	0.626667	0.810000	0.783333	0.190000
2	23915	Grp261	1.130000	0.143333	0.946667	0.123333	0.080000	0.836667	0.056667	0.756667	0.226667
3	50806	Grp198	0.636667	1.090000	0.750000	0.940000	0.743333	0.346667	0.956667	0.633333	0.486667
4	184244	Grp228	0.560000	1.013333	0.593333	0.416667	0.773333	0.460000	0.853333	0.796667	0.516667

```
In [46]: 1 lambdawts['Group'].nunique()
```

```
Out[46]: 1400
```

```
In [47]: 1 lambdawts.shape
```

```
Out[47]: (1400, 2)
```

```
In [48]: 1 all_data.shape
```

```
Out[48]: (284807, 31)
```

```
In [49]: 1 all_data['Group'].nunique()
```

```
Out[49]: 1400
```

```
In [50]: 1 all_data = pd.merge(all_data, lambdawts, on='Group', how='left')
```

```
In [51]: 1 all_data.shape
```

```
Out[51]: (284807, 32)
```

```
In [52]: 1 all_data.head()
```

```
Out[52]:
```

	id	Group	Per1	Per2	Per3	Per4	Per5	Per6	Per7	Per8	Per9
0	112751	Grp169	1.070000	0.580000	0.480000	0.766667	1.233333	1.993333	0.340000	1.010000	0.863333
1	18495	Grp161	0.473333	1.206667	0.883333	1.430000	0.726667	0.626667	0.810000	0.783333	0.190000
2	23915	Grp261	1.130000	0.143333	0.946667	0.123333	0.080000	0.836667	0.056667	0.756667	0.226667
3	50806	Grp198	0.636667	1.090000	0.750000	0.940000	0.743333	0.346667	0.956667	0.633333	0.486667
4	184244	Grp228	0.560000	1.013333	0.593333	0.416667	0.773333	0.460000	0.853333	0.796667	0.516667

```
In [53]: 1 qset['id'].nunique()
```

```
Out[53]: 284807
```

```
In [54]: 1 qset.shape
```

```
Out[54]: (1424035, 2)
```

```
In [55]: 1 qset = qset.groupby('id').mean()
```

```
In [56]: 1 qset.shape
```

```
Out[56]: (284807, 1)
```

```
In [57]: 1 all_data = pd.merge(all_data, qset, on='id', how='left')
```

```
In [58]: 1 all_data.head()
```

```
Out[58]:
```

	id	Group	Per1	Per2	Per3	Per4	Per5	Per6	Per7	Per8	Pe
0	112751	Grp169	1.070000	0.580000	0.480000	0.766667	1.233333	1.993333	0.340000	1.010000	0.8633
1	18495	Grp161	0.473333	1.206667	0.883333	1.430000	0.726667	0.626667	0.810000	0.783333	0.1900
2	23915	Grp261	1.130000	0.143333	0.946667	0.123333	0.080000	0.836667	0.056667	0.756667	0.2266
3	50806	Grp198	0.636667	1.090000	0.750000	0.940000	0.743333	0.346667	0.956667	0.633333	0.4866
4	184244	Grp228	0.560000	1.013333	0.593333	0.416667	0.773333	0.460000	0.853333	0.796667	0.5166

```
In [59]: 1 all_data.tail()
```

```
Out[59]:
```

	id	Group	Per1	Per2	Per3	Per4	Per5	Per6	Per7	Per8	
284802	18333	Grp102	0.553333	1.043333	1.096667	0.686667	0.673333	0.340000	0.900000	0.643333	
284803	244207	Grp504	1.353333	0.616667	0.276667	0.783333	0.690000	0.650000	0.473333	0.670000	
284804	103277	Grp78	1.083333	0.433333	0.806667	0.490000	0.243333	0.316667	0.533333	0.606667	
284805	273294	Grp134	0.566667	1.153333	0.370000	0.616667	0.793333	0.226667	0.910000	0.696667	
284806	223337	Grp18	1.426667	0.110000	-0.006667	-0.200000	0.983333	1.870000	0.033333	0.963333	

```
In [60]: 1 # split the train and test data separately
2 train = all_data[all_data['data']=='train']
3 test = all_data[all_data['data']=='test']
```

```
In [61]: 1 print(train.shape)
2 print(test.shape)

(227845, 33)
(56962, 33)
```

```
In [62]: 1 # Target - train dataset
2 train.head()
```

```
Out[62]:
```

	id	Group	Per1	Per2	Per3	Per4	Per5	Per6	Per7	Per8	Pe
0	112751	Grp169	1.070000	0.580000	0.480000	0.766667	1.233333	1.993333	0.340000	1.010000	0.8633
1	18495	Grp161	0.473333	1.206667	0.883333	1.430000	0.726667	0.626667	0.810000	0.783333	0.1900
2	23915	Grp261	1.130000	0.143333	0.946667	0.123333	0.080000	0.836667	0.056667	0.756667	0.2266
3	50806	Grp198	0.636667	1.090000	0.750000	0.940000	0.743333	0.346667	0.956667	0.633333	0.4866
4	184244	Grp228	0.560000	1.013333	0.593333	0.416667	0.773333	0.460000	0.853333	0.796667	0.5166

```
In [63]: 1 Fraud = train[train['Target']==1]
2 Valid = train[train['Target']==0]
3 outlier_fraction = (len(Fraud))/(len(train))*100
4 print(outlier_fraction)

0.17292457591783889
```

```
In [64]: 1 print(len(Fraud))
2 print(len(Valid))

394
227451
```

```
In [65]: 1 x = train.drop(['id', 'Group', 'Target', 'data'], axis=1)
         2 y = train[['Target']]
```

```
In [66]: 1 x.head()
```

```
Out[66]:
```

	Per1	Per2	Per3	Per4	Per5	Per6	Per7	Per8	Per9	Dem1	
0	1.070000	0.580000	0.480000	0.766667	1.233333	1.993333	0.340000	1.010000	0.863333	0.460000	0.6
1	0.473333	1.206667	0.883333	1.430000	0.726667	0.626667	0.810000	0.783333	0.190000	0.470000	0.6
2	1.130000	0.143333	0.946667	0.123333	0.080000	0.836667	0.056667	0.756667	0.226667	0.660000	0.7
3	0.636667	1.090000	0.750000	0.940000	0.743333	0.346667	0.956667	0.633333	0.486667	1.096667	0.4
4	0.560000	1.013333	0.593333	0.416667	0.773333	0.460000	0.853333	0.796667	0.516667	0.756667	0.6

```
In [67]: 1 y.head()
```

```
Out[67]:
```

	Target
0	0.0
1	0.0
2	0.0
3	0.0
4	0.0

```
In [68]: 1 # Test dataset
         2 test.head()
```

```
Out[68]:
```

	id	Group	Per1	Per2	Per3	Per4	Per5	Per6	Per7	Per8
227845	146574	Grp229	-0.300000	1.540000	0.220000	-0.280000	0.570000	0.260000	0.700000	1.076667
227846	268759	Grp141	0.633333	0.953333	0.810000	0.466667	0.910000	0.253333	1.040000	0.550000
227847	59727	Grp188	1.043333	0.740000	0.860000	1.006667	0.583333	0.616667	0.630000	0.686667
227848	151544	Grp426	1.283333	0.300000	0.576667	0.636667	0.256667	0.543333	0.356667	0.663333
227849	155008	Grp443	1.186667	0.326667	0.476667	0.866667	0.436667	0.680000	0.476667	0.686667

```
In [69]: 1 test = test.drop(['id', 'Group', 'Target', 'data'], axis=1)
```

```
In [70]: 1 test.head()
```

```
Out[70]:
```

	Per1	Per2	Per3	Per4	Per5	Per6	Per7	Per8	Per9	Der
227845	-0.300000	1.540000	0.220000	-0.280000	0.570000	0.260000	0.700000	1.076667	0.930000	0.1566
227846	0.633333	0.953333	0.810000	0.466667	0.910000	0.253333	1.040000	0.550000	0.543333	0.4333
227847	1.043333	0.740000	0.860000	1.006667	0.583333	0.616667	0.630000	0.686667	0.593333	1.2500
227848	1.283333	0.300000	0.576667	0.636667	0.256667	0.543333	0.356667	0.663333	1.156667	1.1866
227849	1.186667	0.326667	0.476667	0.866667	0.436667	0.680000	0.476667	0.686667	1.476667	1.2133

In [71]: 1 x.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 227845 entries, 0 to 227844
Data columns (total 29 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Per1                                227845 non-null  float64
1   Per2                                227845 non-null  float64
2   Per3                                227845 non-null  float64
3   Per4                                227845 non-null  float64
4   Per5                                227845 non-null  float64
5   Per6                                227845 non-null  float64
6   Per7                                227845 non-null  float64
7   Per8                                227845 non-null  float64
8   Per9                                227845 non-null  float64
9   Dem1                                227845 non-null  float64
10  Dem2                                227845 non-null  float64
11  Dem3                                227845 non-null  float64
12  Dem4                                227845 non-null  float64
13  Dem5                                227845 non-null  float64
14  Dem6                                227845 non-null  float64
15  Dem7                                227845 non-null  float64
16  Dem8                                227845 non-null  float64
17  Dem9                                227845 non-null  float64
18  Cred1                               227845 non-null  float64
19  Cred2                               227845 non-null  float64
20  Cred3                               227845 non-null  float64
21  Cred4                               227845 non-null  float64
22  Cred5                               227845 non-null  float64
23  Cred6                               227845 non-null  float64
24  Normalised_FNT                      227845 non-null  float64
25  geo_score                           227845 non-null  float64
26  instance_scores                     227845 non-null  float64
27  lambda_wt                           227845 non-null  float64
28  qsets_normalized_tat                 227845 non-null  float64
dtypes: float64(29)
memory usage: 52.1 MB
```

In [72]: 1 x.describe()

Out[72]:

	Per1	Per2	Per3	Per4	Per5	Per6	
count	227845.000000	227845.000000	227845.000000	227845.000000	227845.000000	227845.000000	227845
mean	0.666006	0.667701	0.666315	0.666687	0.666723	0.667378	0
std	0.654133	0.548305	0.506357	0.471956	0.461393	0.444573	0
min	-18.136667	-23.573333	-15.443333	-1.226667	-37.246667	-8.053333	-13
25%	0.360000	0.470000	0.370000	0.383333	0.436667	0.410000	0
50%	0.670000	0.690000	0.726667	0.660000	0.650000	0.576667	0
75%	1.103333	0.933333	1.010000	0.913333	0.870000	0.800000	0
max	1.483333	8.020000	3.793333	6.163333	12.266667	25.100000	40

```
In [73]: 1 x['Normalised_FNT'].describe()
```

```
Out[73]: count      227845.000000  
mean        -227.954170  
std          61.951661  
min         -250.000000  
25%         -248.617500  
50%         -244.510000  
75%         -230.750000  
max          6172.790000  
Name: Normalised_FNT, dtype: float64
```

```
In [74]: 1 sns.boxplot(y='Normalised_FNT', data=x)  
2 plt.show()
```



```
In [75]: 1 IQR = -230.750000 + 248.617500  
2 IQR
```

```
Out[75]: 17.867500000000007
```

```
In [76]: 1 # pos_outlier_range = Q3 + 1.5*IQR  
2 pos_outlier_range = -230.750000 + (1.5*IQR)  
3 pos_outlier_range
```

```
Out[76]: -203.94875
```

- Holding capping method right now as positive outlier range is -203.94 which is very less and only few data are above this range

Feature scaling

```
In [77]: 1 from sklearn.preprocessing import StandardScaler  
2 sc = StandardScaler()  
3 sc_x = sc.fit_transform(x)
```



```
In [78]: 1 pd.DataFrame(sc_x).describe()
```

```
Out[78]:
```

	0	1	2	3	4	5	
count	2.278450e+05	2.278450e+05	2.278450e+05	2.278450e+05	2.278450e+05	2.278450e+05	2.278450
mean	2.469880e-16	-1.035354e-16	2.169253e-16	3.414173e-16	-1.947214e-16	-2.275907e-16	7.03541
std	1.000002e+00	1.000002e+00	1.000002e+00	1.000002e+00	1.000002e+00	1.000002e+00	1.00000
min	-2.874447e+01	-4.421098e+01	-3.181486e+01	-4.011726e+00	-8.217170e+01	-1.961595e+01	-3.49334
25%	-4.678047e-01	-3.605683e-01	-5.851915e-01	-6.003821e-01	-4.986139e-01	-5.789326e-01	-4.41712
50%	6.105872e-03	4.066939e-02	1.191877e-01	-1.416771e-02	-3.624526e-02	-2.040406e-01	3.14356
75%	6.685615e-01	4.844626e-01	6.787413e-01	5.226069e-01	4.405724e-01	2.983147e-01	4.56467
max	1.249484e+00	1.340918e+01	6.175532e+00	1.164655e+01	2.514117e+01	5.495757e+01	9.67060

Check imbalance dataset

```
In [79]: 1 y.value_counts()
```

```
Out[79]: Target
0.0      227451
1.0        394
dtype: int64
```

```
In [80]: 1 fraud_per = 394/(394+227451)*100
2 fraud_per
```

```
Out[80]: 0.17292457591783889
```

```
In [81]: 1 x.shape
```

```
Out[81]: (227845, 29)
```

```
In [82]: 1 # Since data is imbalance, so we can build model with both aproach
2 # 1) balance the data and perform model building
3 # 2) model building with balance the data
```

```
In [83]: 1 import imblearn
2 from imblearn.over_sampling import SMOTE
3 ros = SMOTE()
4 x_ros, y_ros = ros.fit_resample(sc_x, y)
5 print("before data is imbalance")
6 print(y.value_counts())
7 print()
8 print("after balancing the data: ")
9 print(y_ros.value_counts())
```

before data is imbalance

Target

0.0 227451

1.0 394

dtype: int64

after balancing the data:

Target

0.0 227451

1.0 227451

dtype: int64

Split the data into training and testing for model building

```
In [84]: 1 from sklearn.model_selection import train_test_split
2 x_train, x_test, y_train, y_test = train_test_split(x_ros, y_ros, test_size=0.2)
```

Logistic Regression

```
In [85]: 1 from sklearn.linear_model import LogisticRegression
2 logit = LogisticRegression()
3 logit.fit(x_train, y_train)
```

Out[85]: LogisticRegression()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [86]: 1 y_pred_train = logit.predict(x_train)
2 y_pred_test = logit.predict(x_test)
```

```
In [87]: 1 from sklearn.metrics import confusion_matrix, classification_report, accuracy_score
```

```
In [88]: 1 print(confusion_matrix(y_train, y_pred_train))
2 print()
3 print(confusion_matrix(y_test, y_pred_test))
```

```
[[165950  4677]
 [ 15655 154894]]
```

```
[[55239  1585]
 [ 5258 51644]]
```

```
In [89]: 1 (15477+4647)/(165980+4647+15477+155072)
```

Out[89]: 0.05898421928857833

```
In [90]: 1 print(classification_report(y_train, y_pred_train))
2 print()
3 print(classification_report(y_test, y_pred_test))
```

	precision	recall	f1-score	support
0.0	0.91	0.97	0.94	170627
1.0	0.97	0.91	0.94	170549
accuracy			0.94	341176
macro avg	0.94	0.94	0.94	341176
weighted avg	0.94	0.94	0.94	341176

	precision	recall	f1-score	support
0.0	0.91	0.97	0.94	56824
1.0	0.97	0.91	0.94	56902
accuracy			0.94	113726
macro avg	0.94	0.94	0.94	113726
weighted avg	0.94	0.94	0.94	113726

```
In [91]: 1 print(accuracy_score(y_train, y_pred_train))
2 print()
3 print(accuracy_score(y_test, y_pred_test))
```

0.940406124698103

0.9398290628352356


Decision Tree

```
In [92]: 1 from sklearn.tree import DecisionTreeClassifier
2 dtree= DecisionTreeClassifier(criterion='entropy')
3 dtree.fit(x_train, y_train)
4 y_pred_train_dt = dtree.predict(x_train)
5 y_pred_test_dt = dtree.predict(x_test)
6 print(accuracy_score(y_train, y_pred_train_dt))
7 print()
8 print(accuracy_score(y_test, y_pred_test_dt))
```

1.0

0.9984524207305278

```
In [93]: 1 from sklearn import tree
2 tree.plot_tree(dtrees, filled=True)
3 plt.show()
```



RandomForest Classification

```
In [94]: 1 from sklearn.ensemble import RandomForestClassifier
2 rf = RandomForestClassifier(n_estimators=100,criterion='entropy')
3 rf.fit(x_train, y_train)
4 y_pred_train_rf = rf.predict(x_train)
5 y_pred_test_rf = rf.predict(x_test)
6 print(accuracy_score(y_train, y_pred_train_rf))
7 print()
8 print(accuracy_score(y_test, y_pred_test_rf))
```

1.0

0.9998856901675958

```
In [95]: 1 from xgboost import XGBClassifier
```

XGBoost Classifier

```
In [96]: 1 from xgboost import XGBClassifier
2 xgb = XGBClassifier()
3 xgb.fit(x_train, y_train)
```

```
Out[96]: XGBClassifier(base_score=None, booster=None, callbacks=None,
                      colsample_bylevel=None, colsample_bynode=None,
                      colsample_bytree=None, device=None, early_stopping_rounds=None,
                      enable_categorical=False, eval_metric=None, feature_types=None,
                      gamma=None, grow_policy=None, importance_type=None,
                      interaction_constraints=None, learning_rate=None, max_bin=None,
                      max_cat_threshold=None, max_cat_to_onehot=None,
                      max_delta_step=None, max_depth=None, max_leaves=None,
                      min_child_weight=None, missing=nan, monotone_constraints=None,
                      multi_strategy=None, n_estimators=None, n_jobs=None,
                      num_parallel_tree=None, random_state=None, ...)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [97]: 1 y_pred_train_xg = xgb.predict(x_train)
2 y_pred_test_xg = xgb.predict(x_test)
3 print(accuracy_score(y_train, y_pred_train_xg))
4 print()
5 print(accuracy_score(y_test, y_pred_test_xg))
```

0.9999970689614744

0.9997362080790673

Stacking Classifier

```
In [98]: 1 from mlxtend.classifier import StackingClassifier
2 from sklearn.naive_bayes import GaussianNB
3 from sklearn.model_selection import cross_val_score
```

```
In [99]: 1 clf1 = LogisticRegression()
2 clf2 = DecisionTreeClassifier(criterion='entropy')
3 clf3 = RandomForestClassifier(n_estimators=100, criterion='entropy')
4 clf4 = XGBClassifier()
5 clf5 = GaussianNB()
6 sclf = StackingClassifier(classifiers=[clf2, clf3, clf4, clf5], meta_classifier=
7 print('3-fold cross validation : \n')
8 for clf, label in zip([clf2, clf3, clf4, clf5, sclf], ['Dtree', 'RForest', 'XGBoost', 'Naive_Bayes', 'StackingClassifier']):
9     scores = cross_val_score(clf, x_train, y_train, cv=3, scoring='accuracy')
10    print("Accuracy : %0.2f (+/-%0.2f)[%s]" % (scores.mean(), scores.std(), label))
```

3-fold cross validation :

Accuracy : 1.00 (+/-0.00)[Dtree]
Accuracy : 1.00 (+/-0.00)[RForest]
Accuracy : 1.00 (+/-0.00)[XGBoost]
Accuracy : 0.91 (+/-0.00)[Naive_Bayes]
Accuracy : nan (+/-nan)[StackingClassifier]

```
In [100]: 1 # Anomaly Detection Model
          2 # 1) IsolationForest - RF
          3 # 2) LocalOutlierFactor -knn
          4 # 3) OneClassSVM - SVM
```

```
In [101]: 1 from sklearn.ensemble import IsolationForest
          2 from sklearn.neighbors import LocalOutlierFactor
          3 from sklearn.svm import OneClassSVM
```

```
In [102]: 1 classification = {'IsolationForest' : IsolationForest(contamination=outlier
          2                                     "LocalOutlierFactor": LocalOutlierFactor(contamination=out
          3                                     "OneClassSVM" : OneClassSVM())}
```

```
In [103]: 1 n_outlier = len(Fraud)
          2 n_outlier
```

Out[103]: 394

```
In [104]: 1 fraud_percent=len(Fraud)*100/len(x)
          2 print('Percentage of Fraud :',fraud_percent)
```

Percentage of Fraud : 0.17292457591783889

```
In [105]: 1 from sklearn.metrics import confusion_matrix, classification_report, accuracy_score
```

```
In [106]: 1 x.shape
```

Out[106]: (227845, 29)

```
In [107]: 1 y.shape
```

Out[107]: (227845, 1)

```
In [108]: 1 y.head()
```

Out[108]:

	Target
0	0.0
1	0.0
2	0.0
3	0.0
4	0.0

Anamoly Detection:

```
In [ ]: 1 for i, (clf_name, clf) in enumerate(classification.items()):
2         if clf_name == 'LocalOutlierFactor':
3             y_pred = clf.fit_predict(x)
4             score_prediction = clf.negative_outlier_factor_
5
6         elif clf_name == 'OneClassSVM':
7             clf.fit(x)
8             y_pred = clf.predict(x)
9
10        else:
11            clf.fit(x)
12            score_prediction = clf.decision_function(x)
13            y_pred = clf.predict(x)
14
15        y_pred[y_pred == 1] = 0
16        y_pred[y_pred == -1] = 1
17        n_error = (y_pred != 1).sum()
18
19        print("{} : {}".format(clf_name, n_error))
20        print()
21
22        print("Accuracy Score :")
23        print(accuracy_score(y, y_pred))
24        print()
25
26        print("Classification Report :")
27        print(classification_report(y, y_pred))
```

IsolationForest : 188445

Accuracy Score :
0.8284974434374246

Classification Report :

	precision	recall	f1-score	support
0.0	1.00	0.83	0.91	227451
1.0	0.01	0.91	0.02	394
accuracy			0.83	227845
macro avg	0.50	0.87	0.46	227845
weighted avg	1.00	0.83	0.90	227845

LocalOutlierFactor : 188445

Accuracy Score :
0.8265926397331519

Classification Report :

	precision	recall	f1-score	support
0.0	1.00	0.83	0.91	227451
1.0	0.00	0.36	0.01	394
accuracy			0.83	227845
macro avg	0.50	0.59	0.46	227845
weighted avg	1.00	0.83	0.90	227845

In []:

1